

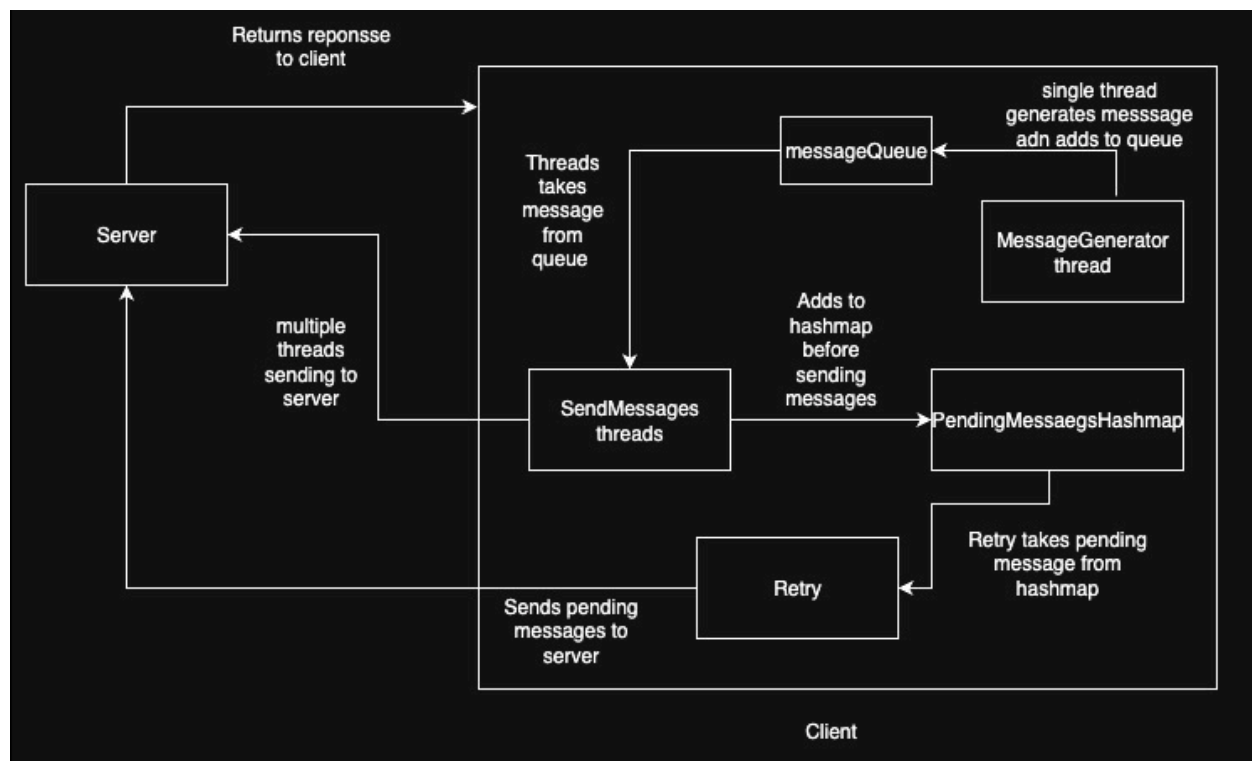
# Assignment 1

## Git repository URL

/server	<a href="https://github.com/RonaldRommel/Chatflow/tree/main/server">https://github.com/RonaldRommel/Chatflow/tree/main/server</a>
/client-part1	<a href="https://github.com/RonaldRommel/Chatflow/tree/main/client-part1">https://github.com/RonaldRommel/Chatflow/tree/main/client-part1</a>
/client-part2	<a href="https://github.com/RonaldRommel/Chatflow/tree/main/client-part2">https://github.com/RonaldRommel/Chatflow/tree/main/client-part2</a>
/results	<a href="https://github.com/RonaldRommel/Chatflow/tree/main/results">https://github.com/RonaldRommel/Chatflow/tree/main/results</a>

## Design

### Architecture Diagram



## Major Classes and their relationships

Class	Responsibility	Relationships
<b>Server</b>		
WebSocketServerApplication	Entry point of the server application. Bootstraps Spring Boot or similar backend and initializes WebSocket configuration.	Instantiates WebSocketConfig and ServerController.
WebSocketHandler	Core handler for WebSocket sessions. Receives messages from clients, parses them, and broadcasts responses. Handles message types (JOIN, CHAT, ACK).	Called by the WebSocket framework when a message/event occurs. Interacts with ServerController.
ServerController	Provides REST endpoints such as /health for readiness/liveness probes and monitoring.	Exposed separately via HTTP.
WebSocketConfig	Configures WebSocket endpoints, message broker, allowed origins, and handlers.	Registers WebSocketHandler. Used by WebSocketServerApplication.
<b>Client</b>		
ClientApplication	Entry point for client execution. Initializes configuration, threads, and test parameters (user count, message rate, etc.).	Creates and manages multiple LoadTestClient instances.
LoadTestClient	Represents a single simulated WebSocket user. Connects to the server, sends ChatMessages, receives ChatResponses. Tracks latency and throughput.	Uses ChatMessage, PendingMessage, BufferedCSVWriter, StreamingLatency. Interacts with server over WebSocket.
ChatMessage	Represents outgoing messages (e.g., join room, send chat text). Contains messageId, roomId, timestamp, and message type.	Created by LoadTestClient and stored in PendingMessage. Serialized and sent via WebSocket.

ChatResponse	Represents a message received from the server (acknowledgment or chat broadcast). Used to measure round-trip latency.	Processed by LoadTestClient. Correlates with PendingMessage via messageId.
PendingMessage	Stores metadata about an in-flight message (send timestamp, retries, etc.) until a response is received.	Managed by LoadTestClient's pending message map. Updated on send/ack.
BufferedCSVWriter	Writes latency and throughput metrics to CSV in batches (buffered mode to avoid I/O overhead).	Used by LoadTestClient or StreamingLatency to persist metrics.
StreamingLatency	Continuously collects latency data (from PendingMessage updates) and computes moving averages or percentiles.	Used by LoadTestClient for live latency reporting; outputs to BufferedCSVWriter.
ThroughputChart	Aggregates message send/ack counts per time window to visualize system throughput over time.	Consumes metrics from StreamingLatency or logs produced by multiple clients.

## Threading Model Explanation

Component	Threading Model	Description
<b>WebSocketHandler</b>	Event-driven, non-blocking I/O	The WebSocket server runs on a reactive thread model. Each connection is handled asynchronously no dedicated thread per connection. Incoming messages are processed in the I/O thread and dispatched to internal handlers via lightweight callbacks.
<b>ServerController (/health)</b>	Separate HTTP thread pool	REST endpoints like <b>/health</b> execute on a dedicated HTTP thread pool (e.g., Tomcat worker threads). These threads are independent of WebSocket I/O threads, ensuring that health checks do not interfere with live message handling.

Thread	Role	Interactions
<b>MessageGenerator Thread (1 per client)</b>	Continuously generates chat messages at a configured rate and enqueues them into a shared messageQueue.	Produces messages for worker threads.
<b>SendMessage Threads (N per client)</b>	Multiple worker threads consume messages from the queue and send them to the WebSocket server. Each message is recorded in the PendingMessagesHashMap along with its timestamp for latency tracking.	Producer-consumer relationship with MessageGenerator.
<b>Retry Thread (1 per client)</b>	Periodically scans the PendingMessagesHashMap to find messages that have not received a server ACK within a timeout. Retries those messages and updates retry metrics.	Ensures reliability and consistency.

## WebSocket Connection Management Strategy

The client manages multiple persistent WebSocket connections to simulate real-time chat users sending and receiving messages concurrently.

### 1) Connection Pool Design

- a) The client pre-creates a pool of WebSocket connections (POOL\_SIZE = 20) during startup.
- b) Each connection is associated with a unique room ID to simulate traffic across different chat rooms.
- c) Connections are stored in a shared list connectionPool, which is used by sender threads to distribute messages evenly.

### 2) Connection Lifecycle

Stage	Description
-------	-------------

Creation	Each connection is instantiated as a WebSocketClient object with custom handlers for onOpen, onMessage, onClose, and onError.
Open	On successful handshake, connectionCount increments and the client becomes ready to send/receive messages.
Active	Connections remain open throughout the test phase, handling both outgoing messages and incoming acknowledgments concurrently.
Close	After all test phases complete, connections are closed gracefully using closeBlocking() to ensure full cleanup.
Error Handling	Failed or broken connections are logged via onError. Metrics track broken connections for reliability analysis.

### 3) Message Routing & Connection Selection

When sending messages:

- Threads pick an active WebSocket connection from the pool in a round-robin manner
- This ensures even load distribution across all open connections.
- If a connection is temporarily closed, the loop skips it and tries the next available one.

## Little's Law Calculation and Predictions

Formula:

$$L = \lambda \times W$$

### Prediction:

RTT for 1 message: 150ms

Suppose each thread can send 100messages/sec.

With 20 connections its 2000 messages

Total messages = 2,000

Test duration  $\approx 0.15$  s  $\rightarrow \lambda \approx 2000 / 0.15 = 13333.33$  messages/s

## Actual:

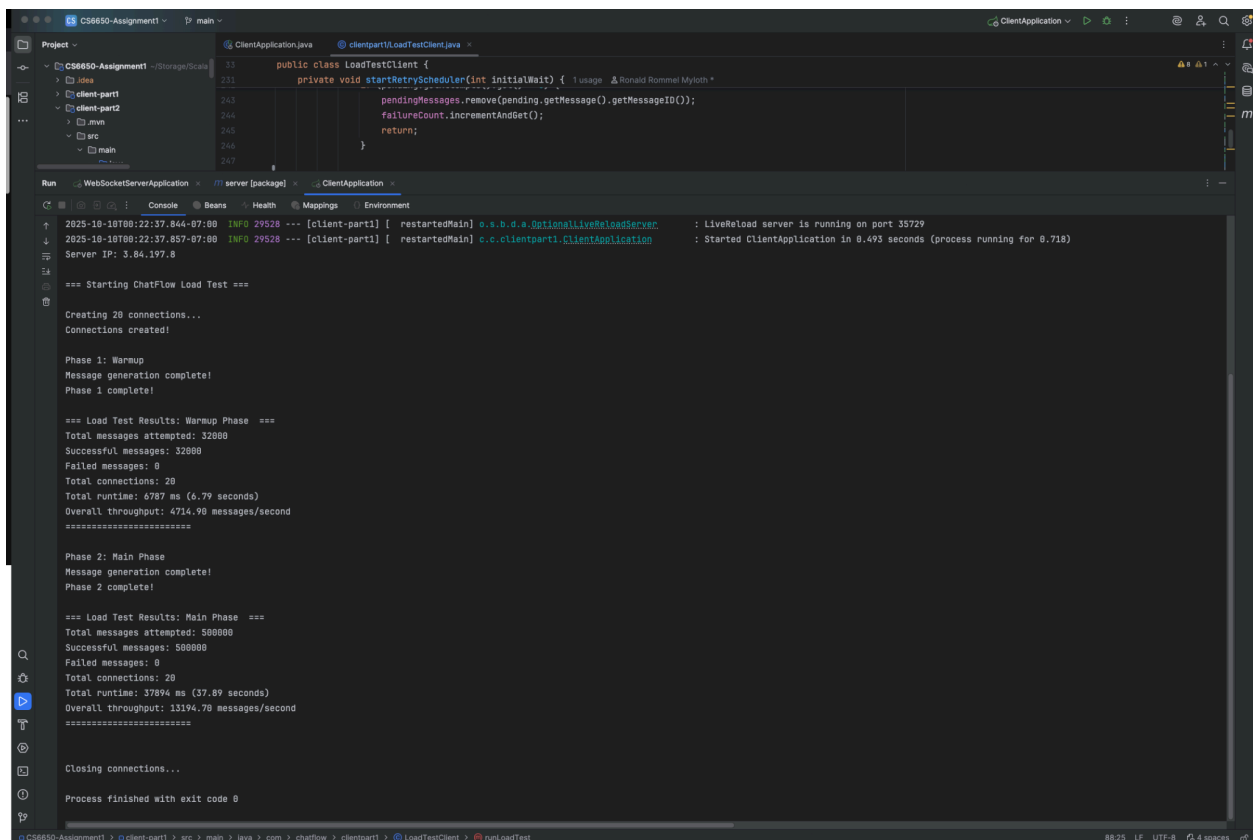
Mean latency (W) = 15,114 ms = 15.114 s

Total messages = 500,000

Test duration  $\approx 37.89$  s  $\rightarrow \lambda \approx 500000 / 37.89 = 13196.10$  messages/s

## Test Results

Screenshot of Part 1 output (basic metrics)



The screenshot shows an IDE with a project named 'CS6650-Assignment1'. The code editor displays the 'LoadTestClient.java' file, which contains a 'LoadTestClient' class with a 'startRetryScheduler' method. The console output shows the execution of the 'LoadTestClient' class, including the 'startRetryScheduler' method call and the 'LoadTestClient' class initialization. The console output also shows the results of the load test, including the number of messages attempted, successful messages, failed messages, total connections, total runtime, and overall throughput.

```
public class LoadTestClient {
    private void startRetryScheduler(int initialWait) {
        pendingMessages.remove(pending.getMessage().getMessageID());
        failureCount.incrementAndGet();
        return;
    }
}

2025-10-10T08:22:37.844-07:00 INFO 29528 --- [client-part1] [ restartedMain] o.s.b.d.a.OptionalLiveReloadServer : LiveReload server is running on port 35729
2025-10-10T08:22:37.857-07:00 INFO 29528 --- [client-part1] [ restartedMain] c.c.clientpart1.ClientApplication : Started ClientApplication in 0.493 seconds (process running for 0.718)
Server IP: 3.84.197.8

=== Starting ChatFlow Load Test ===

Creating 20 connections...
Connections created!

Phase 1: Warmup
Message generation complete!
Phase 1 complete!

=== Load Test Results: Warmup Phase ===
Total messages attempted: 32000
Successful messages: 32000
Failed messages: 0
Total connections: 20
Total runtime: 6787 ms (6.79 seconds)
Overall throughput: 4714.98 messages/second
=====

Phase 2: Main Phase
Message generation complete!
Phase 2 complete!

=== Load Test Results: Main Phase ===
Total messages attempted: 500000
Successful messages: 500000
Failed messages: 0
Total connections: 20
Total runtime: 37894 ms (37.89 seconds)
Overall throughput: 13194.70 messages/second
=====

Closing connections...
Process finished with exit code 0
```

## Screenshot of Part 2 output (detailed metrics)

```
Project > LoadTestClient.java ThroughputChart.java throughput_chart.png latency_metrics.csv ClientApplication2.java >
Run > WebSocketServerApplication > / server [package] > ClientApplication2 >
Console > Beans > Health > Mappings > Environment >
...
=== Load Test Results: Main Phase ===
Total messages attempted: 500000
Successful messages: 500000
Failed messages: 0
Total connections: 20
Total runtime: 35019 ms (35.02 seconds)
Overall throughput: 14277.96 messages/second
=====

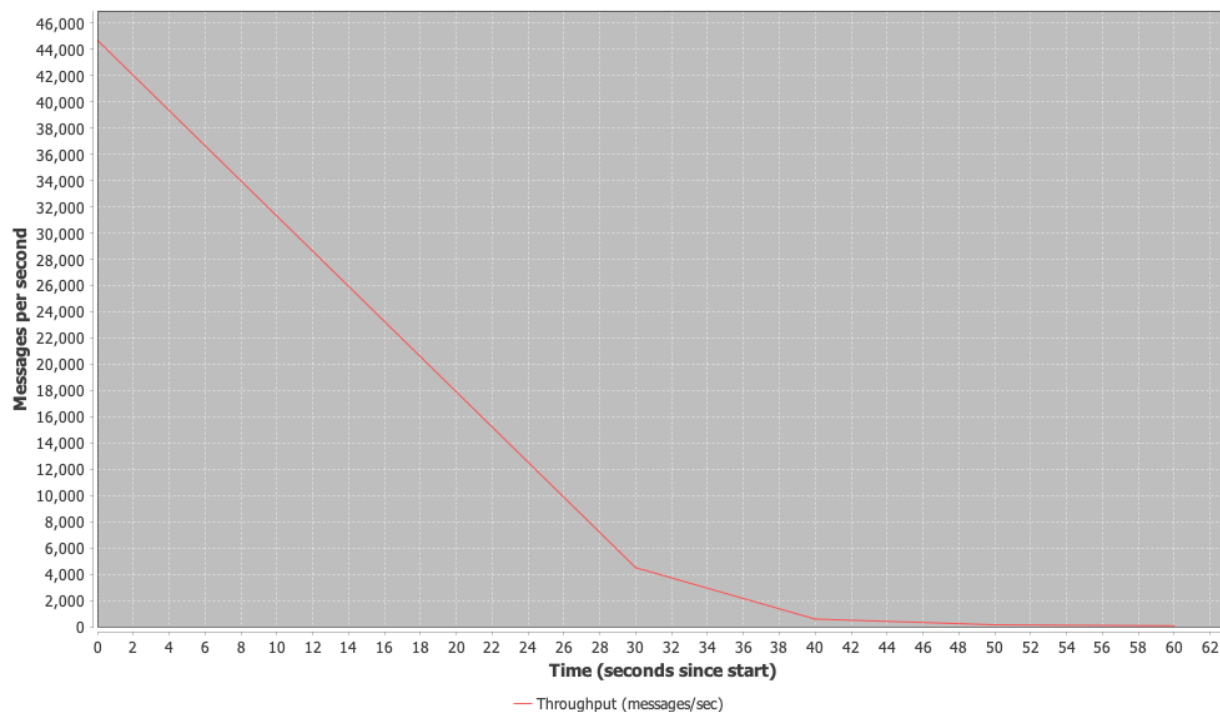
Closing connections...
=== Full Latency Statistics (ms) ===
Count: 500000
Mean response time: 12536.847
msMedian response time: 12253
Min response time: 159 ms
Max response time: 27946 ms
95th percentile response time: 23839
99th percentile response time: 26882

=== Throughput per Room ===
room5: 31250
room6: 31250
room28: 15625
room18: 31250
room3: 31250
room4: 31250
room11: 31250
room9: 31250
room12: 31250
room13: 15625
room7: 31250
room14: 15625
room15: 15625
room8: 31250
room1: 31250
room2: 31250
room16: 15625
room17: 15625
room18: 15625
room19: 15625

=== Message Type Distribution ===
JOIN: 24939 (4.99%)
LEAVE: 25184 (5.02%)
TEXT: 449957 (89.99%)
✓ Throughput chart saved to: ./results/throughput_chart.png
```

## Performance analysis charts

Throughput Over Time



# Evidence of EC2 deployment

```
~Storage/Scalable Systems -- ec2-user@ip-172-31-31-253 ~ -- ssh -i chatflow-key.pem ec2-user@3.84.197.8
~Storage/Scalable Systems/CS6620-Assignment1/server -- ~ssh

(base) ronaldrommelmyloth@Ronald-Macbook Scalable Systems % ssh -i chatflow-key.pem ec2-user@3.84.197.8
#
#####
      _ _ _ _ _
     / _ _ _ _ \
    / _ _ _ _ \
   / _ _ _ _ \
  / _ _ _ _ \
 / _ _ _ _ \
/_ _ _ _ _ \
/m/!
      _ _ _ _ _
     / _ _ _ _ \
    / _ _ _ _ \
   / _ _ _ _ \
  / _ _ _ _ \
 / _ _ _ _ \
/_ _ _ _ _ \

Last login: Fri Oct 10 06:24:51 2025 from 24.19.200.178
[ec2-user@ip-172-31-31-253 ~]$ java -jar server-0.0.1-SNAPSHOT.jar

:: Spring Boot ::
              (v3.5.6)

2025-10-10T07:29:14.218Z INFO 4461 --- [server] [
2-user/server-0.0.1-SNAPSHOT.jar started by ec2-user in /home/ec2-user)
2025-10-10T07:29:14.233Z INFO 4461 --- [server] [
main] c.c.server.WebSocketServerApplication : Starting WebSocketServerApplication v0.0.1-SNAPSHOT using Java 17.0.16 with PID 4461 (/home/ec
2025-10-10T07:29:16.869Z INFO 4461 --- [server] [
main] c.c.server.WebSocketServerApplication : No active profile set, falling back to 1 default profile: "default"
2025-10-10T07:29:16.888Z INFO 4461 --- [server] [
main] o.s.b.w.embedded.tomcat.TomcatWebServer : Tomcat initialized with port 8080 (http)
2025-10-10T07:29:16.889Z INFO 4461 --- [server] [
main] o.apache.catalina.core.StandardService : Starting service [Tomcat]
2025-10-10T07:29:16.892Z INFO 4461 --- [server] [
main] o.apache.catalina.core.StandardEngine : Starting Servlet engine: [Apache Tomcat/10.1.46]
2025-10-10T07:29:16.938Z INFO 4461 --- [server] [
main] o.a.c.c.C.[Tomcat].[localhost].[/] : Initializing Spring embedded WebApplicationContext
2025-10-10T07:29:16.948Z INFO 4461 --- [server] [
main] w.s.c.ServletWebServerApplicationContext : Root WebApplicationContext: initialization completed in 2586 ms
2025-10-10T07:29:17.912Z INFO 4461 --- [server] [
main] o.s.b.a.e.web.EndpointLinksResolver : Exposing 1 endpoint beneath base path '/actuator'
2025-10-10T07:29:18.004Z INFO 4461 --- [server] [
main] o.s.b.w.embedded.tomcat.TomcatWebServer : Tomcat started on port 8080 (http) with context path '/'
2025-10-10T07:29:18.096Z INFO 4461 --- [server] [
main] c.c.server.WebSocketServerApplication : Started WebSocketServerApplication in 4.658 seconds (process running for 5.38)
```